

Your first migration eval

A guided starting point inside Fireworks Evaluations.

[Open the interactive prototype](#) | Desktop experience. No credentials required.

Target user and problem

A developer has a working Claude or GPT prompt and a few examples of correct behavior, but no evaluation setup. Switching models raises a practical question: what will stop working? The hypothesis is that a guided first evaluation reduces setup effort and makes the initial comparison easier to understand. This is an assumption for the take-home, not validated customer research.

Build on existing Fireworks capabilities

Fireworks already provides datasets, evaluation jobs, custom evaluators, and GEPA prompt optimization through Eval Protocol. This concept is an onboarding flow inside Evaluations. Production would turn the prompt and starter examples into a reusable dataset and eval, then expose existing optimization capabilities. Developers with an established eval suite should use that existing workflow.

The critical path

Choose the current provider, bring a prompt, and enter examples with expected outputs. Run a comparison, inspect migration regressions, review suggested prompt additions, and rerun the candidate. Keep current-model, original-candidate, and adapted-candidate results visible together. Export the prompts, examples, and comparison. The demo includes four ready-made invoice examples.

Key choices

Expected values define correctness. The incumbent is evaluated too; shared errors are separate from migration regressions. **Minimal setup.** Enter examples; view model outputs in results. **Start small.** Four examples and four exact field checks illustrate the first evaluation, not production readiness. **Review before applying.** Prompt changes are visible; examples, requirements, and incumbent results stay fixed during adaptation.

What the prototype actually does

Editing, JSON validation, comparison, regression classification, and JSON export work. Both model responses and prompt suggestions use a limited local rule simulation. Claude and OpenAI selections share the same illustrative baseline behavior; Kimi is the sample target. No model calls or credentials are used. The demo resolves three migration regressions and retains one shared currency error. It neither measures model quality nor creates a hosted Fireworks eval.

What I would validate next

Measure time to a first interpretable migration evaluation and setup completion versus the existing path. Check whether users can correctly explain remaining failures. For production, reuse Fireworks datasets, evaluators, and GEPA; connect the incumbent provider or import recorded outputs. Expand starter examples into representative held-out data and repeat runs before making a migration decision.