

Migration check

Identify behavior lost in a model switch, then inspect an adaptation.

[Open the interactive prototype](#) | Desktop experience. No credentials required.

Problem and hypothesis

Changing models can alter application behavior even when the API call still works. For a developer migrating a working extraction workflow from Claude to Fireworks, the question is which requirements stop passing. **Hypothesis:** a guided comparison against their incumbent reduces the effort of diagnosing migration regressions. This workflow gap is assumed, not validated customer research.

Build on what Fireworks already has

Fireworks already provides datasets, evaluation jobs, custom evaluators, and GEPA prompt optimization through Eval Protocol. This concept packages those capabilities into a migration journey in the console. It does not propose a new evaluation or optimization engine; the prototype is not connected to those services.

The workflow

Use a dataset snapshot containing inputs, expected values, and saved Claude outputs. Compare Kimi with the original prompt against the incumbent. Separate **migration regressions** (Claude passes, Kimi fails) from **shared failures** (both fail). Review a proposed adaptation, rerun the same checks, and inspect all three stages together. Export prompts, outputs, and comparison results.

Key choices and assumptions

Expected values define correctness. Claude is scored too; matching it is not sufficient. **Reuse saved outputs.** Assume the user has an incumbent baseline for the same inputs and relevant requirements. **Keep scope narrow.** Four invoice examples, four exact field checks, one target model, desktop only. **Make changes reviewable.** Prompt additions are visible; requirements and baseline stay fixed during adaptation.

Prototype boundaries

Editing, JSON validation, comparison, regression classification, and export work. Claude outputs are authored fixtures; Kimi responses and suggestions use a limited local rule simulation. No model calls or credentials are used. The sample goes from three migration regressions to zero, with one shared currency failure remaining. These invented examples are neither measured model performance nor evidence of migration readiness.

What I would validate next

Can this guided workflow reduce time to a justified migration decision compared with the existing eval and GEPA workflow? A production version would reuse existing datasets and evaluators, preserve model and prompt versions, and check revisions on representative held-out data with repeated runs. The current demo illustrates that proposed experience; it does not establish demand or migration success.

References: [Fireworks Eval Protocol + GEPA](#) | [Custom evaluators](#) | [Prompt migration case study](#)