

Your first migration eval

A guided starting point inside Fireworks Evaluations.

[Open the interactive prototype](#) | Desktop experience. No credentials required.

Target user and problem

A developer has a working Claude or GPT extraction prompt and a few examples of correct behavior, but no evaluation setup. Switching models raises a practical question: what will stop working? The hypothesis is that a guided first evaluation reduces setup effort and makes the initial comparison easier to understand. This is an assumption for the take-home, not validated customer research.

Build on existing Fireworks capabilities

Fireworks already provides datasets, evaluation jobs, custom evaluators, and GEPA prompt optimization through Eval Protocol. This concept is an onboarding flow inside Evaluations. Production would turn the prompt and starter examples into a reusable dataset and eval, then expose existing optimization capabilities. Developers with an established eval suite should use that existing workflow.

The critical path

Three screens keep the task focused: set up a comparison, inspect results, then choose prompt edits. Select current and candidate models; bring a prompt and examples with expected outputs. Inspect individual failures, select only the additions you want, and rerun. Compare the current model, original candidate, and adapted candidate; copy the tested prompt. The demo has four editable invoice examples and a fixed four-field JSON schema.

Key choices

Expected values define correctness. The incumbent is evaluated too; shared errors are separate from migration regressions. **Minimal setup.** Enter examples; view model outputs in results. **Scope explicitly.** Invoice extraction uses four exact field checks; report both field and complete-example passes. **Selective adaptation.** Check each addition and preview the resulting prompt. Only selected edits are tested; examples, requirements, and incumbent results stay fixed.

What the prototype actually does

Editing, JSON validation, comparison, regression classification, and prompt copying work. Both model responses and prompt suggestions use a limited local rule simulation. Seven Claude/GPT choices share one baseline profile; eight Fireworks choices share one candidate profile. No model calls or credentials are used. After all three edits, 15/16 fields but only 3/4 complete examples pass; a shared currency error remains. It neither measures model quality nor creates a hosted Fireworks eval.

What I would validate next

Measure time to a first interpretable migration evaluation and setup completion versus the existing path. Check whether users can correctly explain remaining failures. For production, reuse Fireworks datasets, evaluators, and GEPA; connect the incumbent provider or import recorded outputs. Expand starter examples into representative held-out data and repeat runs before making a migration decision.